

Computer Science 62

Lab 2

Wednesday February 3, 2016

In this laboratory, we will use our **Stopwatch** class to measure the efficiency of the **Vector** class. Specifically, we want to see how execution speed is affected by the **increment** parameter. Recall that **increment** is the amount by which the underlying data array is lengthened when the vector requires more space. If **increment** is set to zero, then the size of the data array is doubled. We'll be using the **Vector** class since the **ArrayList** class only doubles the array length and does not give you incremental building as an option. Before you start coding make sure you look over the documentation for the **Vector** class.

1. Begin by closing all of your open Eclipse projects, so that errors in them will not affect your work today. Use **Projects/Close**.
2. Create a new Eclipse project named Lab2. Remember to continue to the window in which you can add the BAILEY variable. Next, copy the file `/common/cs/cs062/labs/lab02/StopWatch.java` into the **src** directory in your new Eclipse project and select **File/Refresh**.
3. Create a new class **VectorTimer**. This class will contain only a **main** method and a few other static methods:

- `public static long run(int maxSize, int increment)`

The **run** method creates a new empty vector of type `Vector<String>` with the specified increment, using `new Vector<String>(0, increment)`. It returns the time that it takes to add **maxSize** strings to the Vector. Use the `Vector<String>` method **add**, and always add the same constant string—your name, for example. To attempt to minimize the impact from garbage collection add the line: `System.gc();` in your run method right before you start the timer.

- `public static ArrayList<Long> trial(int size,
ArrayList<Integer> incrs)`

The **trial** method compares the results from **run** for a fixed size and varying increments. It makes one call to **run** for each entry in the **incrs** vector. The results are returned in an `ArrayList` whose size is the same as that of **incrs**.

- `public static void main(String[] args)`

The `main` method runs several trials and prints the results. Start with increments of 1, 10, and 0; and sizes of 0, 5000, 10000, 15000, You may want to adjust the sizes when you see the results. *Don't forget that Java uses just-in-time compilation so you'll need to first run several trials and discard the results.*

4. Present the output in a table like the one below; see the tutorial below about formatting. The nanosecond precision of `Stopwatch` is too fine; you will need to adjust the scale of the timing values as they are printed, which can vary from computer to computer.

size	linear (1)	linear (10)	double
0	0	0	0
5000	148	14	1
10000	580	58	0
15000	1321	132	0
20000	2733	267	1
25000	4863	491	1
30000	7781	781	1

We will discuss the significance of your results, and those of your classmates, as they appear. Some things to think about: what is the running time (i.e. Big-O running time) of increment vs. double? Does your data accurately reflect this?

More fun...

Once you've got all this working, if you have time we can try out a few additional things:

- What happens with other increments (besides 1 and 10)? Can you predict what the results will look like, for example what do you think a column headed `linear (100)` would look like?
- Rather than just running one experiment per setting, you can run multiple experiments (say 5 or 10) and average the results in your run method. This will be a bit slower, but should give you more accurate results.
- It may be interesting to compare the performance difference between `ArrayList` and `Vector`. `ArrayList` does **NOT** allow you to adjust the increment size; it always doubles the size. However, you can compare the performance of `Vector` vs. `ArrayList` for doubling sizes. Which is faster?

A note on formatting textual output.

The object `System.out` has type `PrintStream`, which in turn has a method `format`. `format` is very general and makes it easy to print the lines in the table. The call

```
System.out.format("First: %8d, second: %-12s%n", num, str);
```

creates a string and prints it. The string is formed by

- replacing `%8d` with the numerical value of `num`, right justified in a field eight characters wide, and
- replacing `%-12s` with the string representation of `str`, left justified in a field twelve characters wide.

If `num` and `str` are 47 and XLVII respectively, then

```
First:      47, second: XLVII
```

is the result of the method call above.

The letters after the percent sign, `d` and `s` in this example, indicate the kind of data being formatted; they are not variables. The sequence `%n` is the OS independent newline character. You may have as many `%` expressions in the format string as you want; they are matched with the arguments that follow. There are *many* more options for format strings; see the Java documentation for the classes `PrintStream` and `Formatter` or the tutorial at:

```
http://java.sun.com/docs/books/tutorial/java/data/numberformat.html
```

for more information.